

# LibreS3: design, challenges, and steps toward reusable libraries

Török Edwin [edwin@skylable.com](mailto:edwin@skylable.com)

Skylable Ltd.

## Concepts

LibreS3's architecture is not novel: it is based on existing, well-established concepts. What is interesting is how these all work together in a real application.

## Monad with exception handling

```
type +α t ( * the type of deferred computations * )
val return : α → α t ( * immediate value to deferred value * )
val ( >=> ) : α t → (α → β t) → β t ( * chain * )
val fail : exn → α t ( * exception to deferred value * )
val try_bind : (unit → α t) → (α → β t) → (exn → β t) → β
t ( * [try_bind m f g] either f or g * )
```


Implementable by monadic concurrency libraries: Lwt, Async, but also by the Result monad.

## Object storage HTTP(S) REST API

**Amazon S3 server** proprietary, runs on Amazon's hardware

**SX server** FOSS cluster storage backend: supports deduplication, no vendor lock-in

**LibreS3** FOSS S3-compatible server: Allows to connect existing S3 clients to an SX cluster.



- an HTTP(S) server that implements the S3 REST API on one side
- an SX client on the other side
- both LibreS3 and SX runs on your own hardware

## any-http

Somewhat similar to how Mirage applications are built

```
module MyApp(H : Httpintf.S) = struct ( * ... * )
let uri = Uri.make ~scheme:"http" ~host ~port ~path () in
H.Client.call meth uri (H.Headers.init ()) (H.Body.empty)
>=> fun (status, headers, bodystream) →
H.Body.to_string bodystream >=> fun body →
( * ... * )
module App = MyApp(Httpservice_cohttp_async)
module App = MyApp(Httpservice_ocsigenserver)
module App = MyApp(Httpservice_cohttp_lwt.Make
  (Cohttp_lwt_unix.Server) (Cohttp_lwt_unix.Client))
```

## Work in progress

**any-io** abstract over Unix vs Lwt\_unix (Async and Mirage too in the future)

**ioconvcomb** combinators for describing Json/XML to OCaml type conversion

Ideas from: “Type-safe functional formatted IO” (Oleg Kiselyov), `cconv`, `meta_conv` and `dyntype`. Benefits: don't have to choose between Yojson/Jsonm, unified tree/streaming interface and error reporting.

## Issues

**double-release of runtime lock** OCamlNet < 3.7.4, ocaml-ssl 0.4.6a tracked down using a debug patch for OCaml's runtime and submitted patch

**Unix.fork in Ocsigenserver** must use `Lwt_unix.fork` otherwise it hangs if an error message is printed, patched upstream

**ocamlbuild SIGPIPE on \*BSD** worked around in 4.02

**Lwt readdir crash** fixed in 2.4.2

**OCamlnet SSL persistent connections** fixed in 3.7.4

**PUT/DELETE support in Ocsigenserver** implemented in 2.4.0

**Ocsigenserver default worker threads** 1000 too much, using 64

**Mirage XSA-90 / CVE-2014-2580** (crashed Dom0 kernel), not used yet

## Libraries

Reusable components of LibreS3, parametrized by a monad: it is not tied to one particular implementation, or to the use of a monadic concurrency library. See also Cohttp, Mirage, ...

## any-cache: Least Recently Used cache with 2 Queues.

Implemented as a result monad functor

```
module Make(R:Result) : sig
type ('ok, 'err) t
val create : int → ('ok, 'err) t
val get : ('ok, 'err) t → notfound:'err → string → ('ok, 'err) R.t
val set : ('ok, 'err) t → string → ('ok, 'err) R.t → unit
val lookup : ('ok, 'err) t → string →
  (string → ('ok, 'err) R.t) → ('ok, 'err) R.t
end
```

monad transformer provided internally

```
let lift f v = M.try_bind (fun () → f v) return fail
```

usable with monadic concurrency libraries

```
val lookup_exn: (α, exn) t → string → (string → α M.t) → α M.t
```

Testing (same code for Lwt, Async and direct mode)

```
module Make(Monad: MonadTest) = struct
( * * )
  Monad.run (lookup cache key1 (compute_data_direct cnt))
```

## Packaging

If user has OPAM it is easy, but ...

User may not have “new enough OCaml” RHEL/CentOS 6 rebuilt .spec file from Fedora on CentOS6 and provide binary rpm

**Libraries requiring newer versions** 3.11.2 3.12.1 4.00.1

“portable” Linux binaries 4.01.0+lsb and 4.01.0+musl+static opam switches if you really need them

**Full source tarball** embedded 3rdparty libraries custom build script, but very slow

**Automation** lack of a full OASIS/OPAM workflow

**Existing** oasis2deb, oasis2opam

**Missing** (oasis|opam)2rpm, (oasis|opam)2bsd


**Wishlist** opam2release (to build full source tarball, .deb, .rpm and \*BSD)

## Additional materials

Paper, bibliography, sample code



<http://goo.gl/jmF0cn>



<http://www.skylable.com/tag/libres3/>